

Shawn Berg
bergsha@oregonstate.edu
May 30th, 2019
CS 475

Project 6 CUDA Monte Carlo Simulation Writeup

Machine this was run on:

This was run on a Google Cloud Compute Instance with the following GPU specifications:

```
Device 0: "Tesla P100-PCIE-16GB"
CUDA Driver Version / Runtime Version      10.1 / 10.1
CUDA Capability Major/Minor version number: 6.0
Total amount of global memory:              16281 MBytes (17071734784 bytes)
(56) Multiprocessors, ( 64) CUDA Cores/MP: 3584 CUDA Cores
GPU Max Clock rate:                        1329 Mhz (1.33 GHz)
Memory Clock rate:                         715 Mhz
Memory Bus Width:                          4096-bit
L2 Cache Size:                             4194304 bytes
Maximum Texture Dimension Size (x,y,z)      1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
Total amount of constant memory:             65536 bytes
Total amount of shared memory per block:     49152 bytes
Total number of registers available per block: 65536
Warp size:                                  32
Maximum number of threads per multiprocessor: 2048
Maximum number of threads per block:         1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
Maximum memory pitch:                       2147483647 bytes
Texture alignment:                           512 bytes
Concurrent copy and kernel execution:        Yes with 2 copy engine(s)
Run time limit on kernels:                   No
Integrated GPU sharing Host Memory:          No
Support host page-locked memory mapping:     Yes
Alignment requirement for Surfaces:          Yes
Device has ECC support:                      Enabled
Device supports Unified Addressing (UVA):     Yes
Device supports Compute Preemption:          Yes
Supports Cooperative Kernel Launch:          Yes
Supports MultiDevice Co-op Kernel Launch:    Yes
Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 4
Compute Mode:
< Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
```

Table (Performance as a function of Threads vs Number of Trials)

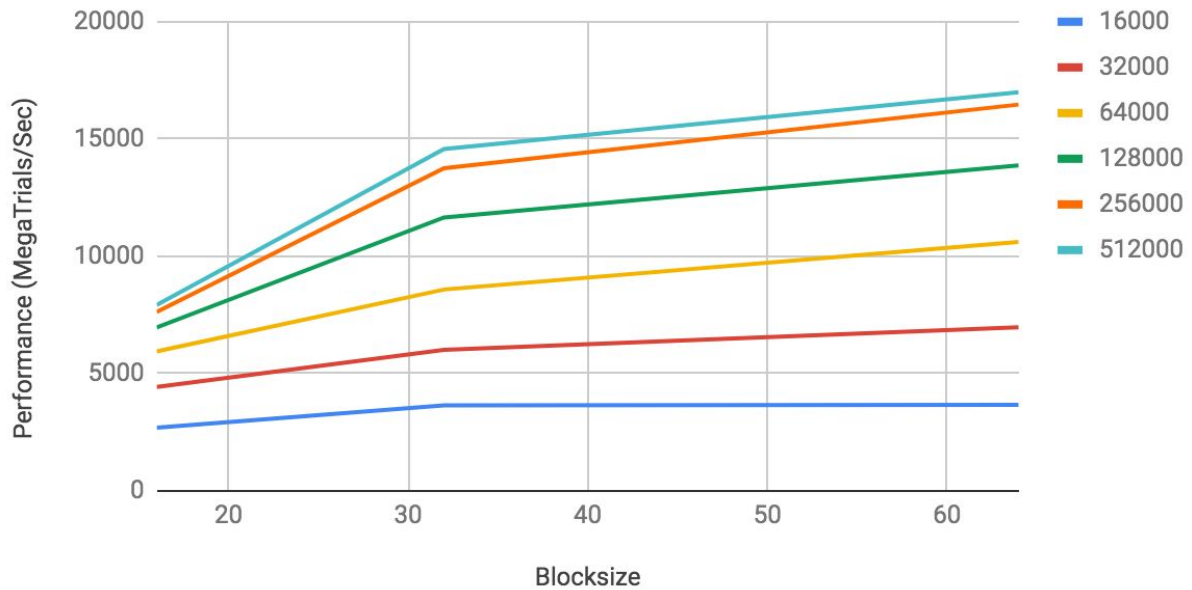
	16000	32000	64000	128000	256000	512000
16	2698.04	4426.93	5949.37	6964.52	7631.48	7926.13
32	3646.71	6009.25	8584.43	11654.67	13758.47	14576.86
64	3669.19	6975.45	10617.4	13873.47	16469.38	16990.55

Units = MegaTrials/Sec

Graph 1 (MegaTrials per Second vs Blocksize)

Monte Carlo Performance

as a function of block size vs array size

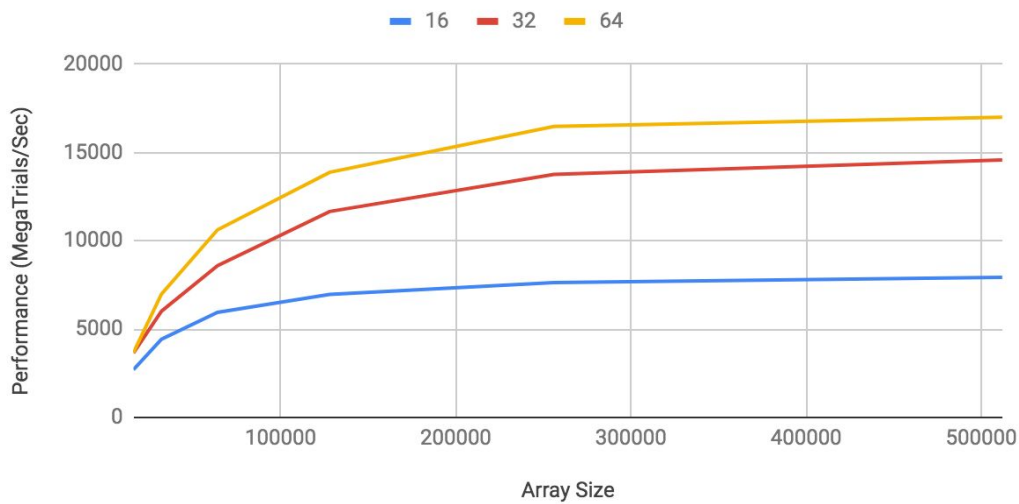


Each line represents array size

Graph 2 (MegaTrials per Second vs Array Size)

Monte Carlo Performance

as a function of array size vs block size



Each line represents blocksize.

Additional Commentary

1. What patterns are you seeing in the performance curves?

The performance curves for both graphs show a general increase in performance as array size and blocksize increases.

2. Why do you think the patterns look this way?

The patterns appear this way because as array size and blocksize increase, more of the GPU is being utilized in the computations, thus an increase in performance is displayed.

3. Why is a BLOCKSIZE of 16 so much worse than the other two?

The blocksize of 16 is worse than the other two because the warp size is 32. This means when using a block size of 16, 32 threads are still allocated to the problem, but only 16 are being used, leaving the other 16 unused.

4. What does that mean for the proper use of GPU parallel computing?

To properly utilize the GPU in parallel computing, not only does the problem size need to be large enough, but the intended division of labor must also take into account the relevant sizes of the blocks on the target device to acquire all of the potential performance.